

DecisionTrees

April 7, 2020

```
[1]: import numpy as np
import pandas as pd
import pydotplus
import matplotlib.pyplot as plt
import matplotlib.patches as patches
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.preprocessing import scale
from sklearn.preprocessing import StandardScaler
from sklearn import tree
from sklearn.tree import DecisionTreeClassifier
from graphviz import Source
import sklearn
from sklearn.model_selection import train_test_split
from IPython.display import Image
import sqlite3
from tqdm import tqdm
import warnings
warnings.filterwarnings('ignore')
warnings.filterwarnings('ignore', 'Solver terminated early.*')
import string
from scipy import interp
from sklearn.model_selection import cross_val_score
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import metrics
from sklearn.metrics import confusion_matrix
from sklearn.calibration import CalibratedClassifierCV , calibration_curve
from sklearn.metrics import f1_score
from sklearn.metrics import roc_curve, auc
import pickle
%matplotlib inline
```

```
[2]: con = sqlite3.connect("/home/niranjan/Downloads/database.sqlite")
data = pd.read_sql_query("select * from Reviews where Score!=3",con)
data['Score'] = [1 if i>3 else 0 for i in data['Score']]
```

Removing Duplicate Data

```
[3]: df = data.sort_values(by= 'Time',ascending=True,inplace=False,kind='quicksort')
data_without_dup = df.
    ↳drop_duplicates(subset={'UserId','ProfileName','Time','Text'},
    ↳inplace=False,keep='first')
data_without_dup = data_without_dup[data_without_dup.
    ↳HelpfulnessNumerator<=data_without_dup.HelpfulnessDenominator]
df_x = data_without_dup.drop(['Score'],axis=1)
```

```
[4]: import nltk
from nltk.corpus import stopwords
from nltk import WordNetLemmatizer
nltk.download('stopwords')
nltk.download('wordnet')
lis = list(stopwords.words('english'))
lem = WordNetLemmatizer()
```

```
[nltk_data] Downloading package stopwords to
[nltk_data] /home/niranjana/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
[nltk_data] Downloading package wordnet to /home/niranjana/nltk_data...
[nltk_data] Package wordnet is already up-to-date!
```

```
[5]: import re
def clean_html(words):
    tag = re.compile(r'<.*?>')
    cleanSent = re.sub(tag, '', words)
    return cleanSent

def punch_remove(words):
    tag = re.compile(r'^a-zA-Z')
    cleanSent = re.sub(tag, '', words)
    return cleanSent
```

****-data cleaning-****

removal of html tags, symbols other than alphabets, stopwords and performing lemmatization as part of data pre-processing.

Lemmatization :- is the process of grouping together the inflected forms of a word so they can be analysed as a single item, identified by the word's lemma, or dictionary form.**

```
[6]: final_data = []
str1 = ' '
```

```

positive_words = []
negative_words = []
i=0
for sen in data_without_dup['Text'].values:
    filter_word = []
    pos_word = []
    neg_word = []
    sent= clean_html(sen)
    for word in sent.split():
        cleanwords = punch_remove(word)
        for cleanword in cleanwords.split():
            if((len(cleanword) >2) & (cleanword.isalpha())):
                if((cleanword.lower() not in lis)):
                    w = (lem.lemmatize(cleanword.lower())).encode('utf-8')
                    filter_word.append(w)
                    if data_without_dup['Score'].values[i] == 1 :
                        pos_word.append(w)
                    else :
                        neg_word.append(w)
                else :
                    continue
            else :
                continue
    str1 = b" ".join(filter_word)
    str2 = b" ".join(pos_word)
    str3 = b" ".join(neg_word)
    final_data.append(str1)
    positive_words.append(str2)
    negative_words.append(str3)
    i = i + 1

```

```

[7]: data_without_dup['final_string'] = final_data
data_without_dup['Positive_string'] = positive_words
data_without_dup['Negative_string'] = negative_words
data_without_dup['final_string'] = data_without_dup['final_string'].str.
    ↳ decode('utf8')
data_without_dup['Positive_string'] = data_without_dup['Positive_string'].str.
    ↳ decode('utf8')
data_without_dup['Negative_string'] = data_without_dup['Negative_string'].str.
    ↳ decode('utf8')

```

```

[8]: X = data_without_dup['final_string']
y = data_without_dup['Score']

```

```

[9]: X_train= X[0:250000]
y_train = y[0:250000]
X_test= X[250000:280000]

```

```
y_test = y[250000:280000]
```

BOW as vectorizer with StandardScaler

```
[12]: count_vect = CountVectorizer()
X_train_bow = count_vect.fit_transform(X_train)
X_test_bow = count_vect.transform(X_test)
count_vec = StandardScaler(with_mean=False)
X_train_bow = count_vec.fit_transform(X_train_bow)
X_test_bow = count_vec.transform(X_test_bow)
```

```
/home/niranjan/anaconda3/lib/python3.6/site-
packages/sklearn/utils/validation.py:595: DataConversionWarning: Data with input
dtype int64 was converted to float64 by StandardScaler.
warnings.warn(msg, DataConversionWarning)
/home/niranjan/anaconda3/lib/python3.6/site-
packages/sklearn/utils/validation.py:595: DataConversionWarning: Data with input
dtype int64 was converted to float64 by StandardScaler.
warnings.warn(msg, DataConversionWarning)
/home/niranjan/anaconda3/lib/python3.6/site-
packages/sklearn/utils/validation.py:595: DataConversionWarning: Data with input
dtype int64 was converted to float64 by StandardScaler.
warnings.warn(msg, DataConversionWarning)
```

Hyper parameter tuning


```
[11]: depth = [1, 5, 10, 50, 100, 500, 100]
min_samples_spl = [5, 10, 100, 500]
fig1 = plt.figure(figsize=[12,12])
ax1 = fig1.add_subplot(111,aspect = 'equal')
ax1.add_patch(
    patches.Arrow(0.45,0.5,-0.25,0.25,width=0.3,color='green',alpha = 0.5)
)
ax1.add_patch(
    patches.Arrow(0.5,0.45,0.25,-0.25,width=0.3,color='red',alpha = 0.5)
)

mean_fpr = np.linspace(0,1,100)
for i in depth:
    for j in min_samples_spl:
        classifier = DecisionTreeClassifier(max_depth=i, min_samples_split=j,
        ↪class_weight='balanced')
        model = CalibratedClassifierCV(classifier,cv=5,method = 'isotonic')
        model.fit(X_train_bow,y_train)
        mod_probs = model.predict_proba(X_test_bow)[:,-1]
        fpr, tpr, thresholds = metrics.roc_curve(y_test, mod_probs)
        auc = metrics.roc_auc_score(y_test, mod_probs)
```

```

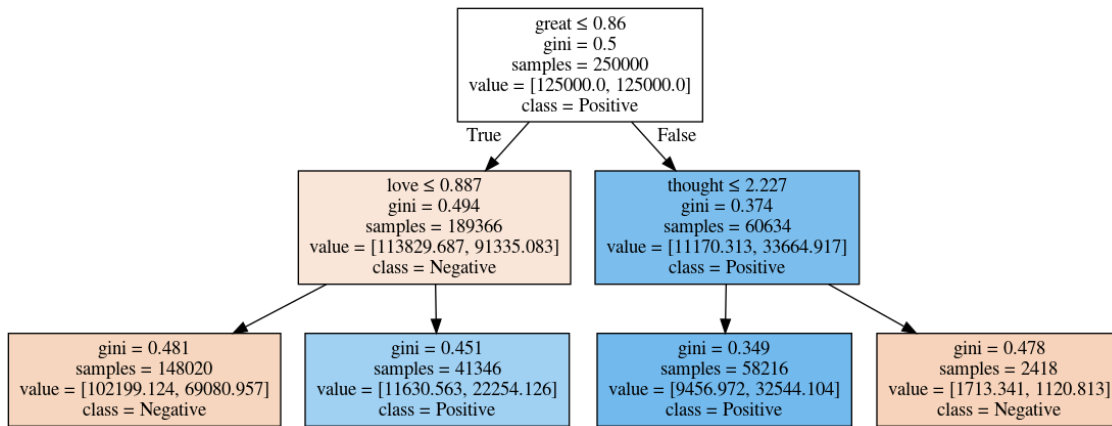
plt.plot(fpr, tpr, label="auc_score/max_depth/min_samples_split_
↪ "+str(auc) + "/" + str(i) + "/" + str(j))
plt.legend(loc=4)
plt.plot([0,1],[0,1], linestyle = '--', lw = 2, color = 'black')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC')
plt.legend(loc="lower right")
plt.text(0.32, 0.7, 'More accurate area', fontsize = 12)
plt.text(0.63, 0.4, 'Less accurate area', fontsize = 12)
plt.show()

```




```
[32]: import graphviz
import pydotplus
from IPython.display import Image
classifier = DecisionTreeClassifier(max_depth=2, class_weight='balanced')
classifier.fit(X_train_bow, y_train)
feature_name = count_vect.get_feature_names()
target = ['Negative', 'Positive']
from sklearn.tree import export_graphviz
graph = tree.export_graphviz(classifier, out_file=None, class_names=target,
    ↪, feature_names = feature_name, filled = True, special_characters=True)
# Draw graph
graph = pydotplus.graph_from_dot_data(graph)
# Show graph
Image(graph.create_png())
```

[32]:




```
[40]: classifier = DecisionTreeClassifier(max_depth=15, class_weight='balanced')
classifier.fit(X_train_bow, y_train)
#coef = classifier.coef_
probs = classifier.predict_proba(X_test_bow)[: , 1]
model = CalibratedClassifierCV(classifier, cv=5, method='isotonic')
model.fit(X_train_bow, y_train)
mod_probs = model.predict_proba(X_test_bow)[: , 1]
#reliability diagram
fop, mpv = calibration_curve(y_test, probs, n_bins=10, normalize=True)
fop1, mpv1 = calibration_curve(y_test, mod_probs, n_bins=10, normalize=True)
# plot perfectly calibrated
plt.plot([0, 1], [0, 1], linestyle='--', label='perfectly calibrated')
# plot model reliability
plt.plot(mpv, fop, marker='.', label='model uncalibrated')
plt.plot(mpv1, fop1, marker='.', label='model calibrated')
plt.title("Reliability Diagram")
plt.xlabel("mean_predicted_values")
```

```
plt.ylabel("fraction_of_positives")
plt.legend()
plt.show()
```



Reliability Diagram : Observed frequency of an event plotted against the Forecast probability of an event.

```
[13]: classifier = DecisionTreeClassifier(max_depth=100, min_samples_split=500,
    ↳class_weight='balanced')
classifier.fit(X_train_bow,y_train)
model = CalibratedClassifierCV(classifier,cv=5,method = 'isotonic')
model.fit(X_train_bow,y_train)
y_pred = model.predict(X_test_bow)
acc = f1_score(y_pred,y_test,average='micro')*100
df = pd.DataFrame(confusion_matrix(y_test,y_pred))
sns.heatmap(df,annot=True,fmt="d")
plt.show()
print('\n\nThe accuracy of the DecisionTreeClassifier for depth = %d and
    ↳min_samples_split = %d is %f%%' % (100,500, acc))
```



The accuracy of the DecisionTreeClassifier for depth = 100 and min_samples_split = 500 is 85.700000%

```
[14]: y_pred = model.predict(X_test_bow)
from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.80	0.26	0.40	5346
1	0.86	0.99	0.92	24654
micro avg	0.86	0.86	0.86	30000
macro avg	0.83	0.62	0.66	30000
weighted avg	0.85	0.86	0.83	30000

```
[15]: coef = classifier.feature_importances_
class_labels = model.classes_
feature_names = count_vect.get_feature_names()
topn_class1 = sorted(zip(coef, feature_names), reverse=False)[:20]
topn_class2 = sorted(zip(coef, feature_names), reverse=True)[:20]
print("Important words in negative reviews")
for coef, feat in topn_class1:
    print(class_labels[0], coef, feat)
```



```

print("-----")
print("Important words in positive reviews")
for coef, feat in topn_class2:
    print(class_labels[1], coef, feat)

```

Important words in negative reviews

```

0 0.0 aa
0 0.0 aaa
0 0.0 aaaa
0 0.0 aaaaa
0 0.0 aaaaaaaaa
0 0.0 aaaaaaaaaaaaaa
0 0.0 aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaabr
0 0.0 aaaaaaaaaaaaaaaaaaaaaargh
0 0.0 aaaaaaaaagghh
0 0.0 aaaaaaahhhhhh
0 0.0 aaaaaarrrrrrggghh
0 0.0 aaaaaabr
0 0.0 aaaaaah
0 0.0 aaaaaahhh
0 0.0 aaaaaahhhhbr
0 0.0 aaaaaahhhhhyaaaaaa
0 0.0 aaaaaawwwwwwww
0 0.0 aaaaaah
0 0.0 aaaaahhhhhhhhhhhhhhhthe
0 0.0 aaaaawill

```

Important words in positive reviews

```

1 0.0990350154071816 great
1 0.05358931950541775 love
1 0.04721148244145879 best
1 0.036247023890801 delicious
1 0.03024856779790211 disappointed
1 0.026546502488069848 bad
1 0.022347812312398498 perfect
1 0.020542713887296185 thought
1 0.019491600028405522 good
1 0.019195535082647958 favorite
1 0.018327058204032076 excellent
1 0.016600775553961878 would
1 0.014625911816804999 money
1 0.01213348699346732 worst
1 0.011647719525502632 wonderful
1 0.011115672121797858 awful
1 0.010046053897243981 highly
1 0.009916758070424038 review
1 0.009036311856948726 unfortunately

```


1 0.00901390900928107 nice

tfidf as vectorizer with StandardScaler

```
[127]: tfidf_vec = TfidfVectorizer()
X_train_tfidf = tfidf_vec.fit_transform(X_train)
X_test_tfidf = tfidf_vec.transform(X_test)
vec = StandardScaler(with_mean=False)
X_train_tfidf = vec.fit_transform(X_train_tfidf)
X_test_tfidf = vec.transform(X_test_tfidf)
```

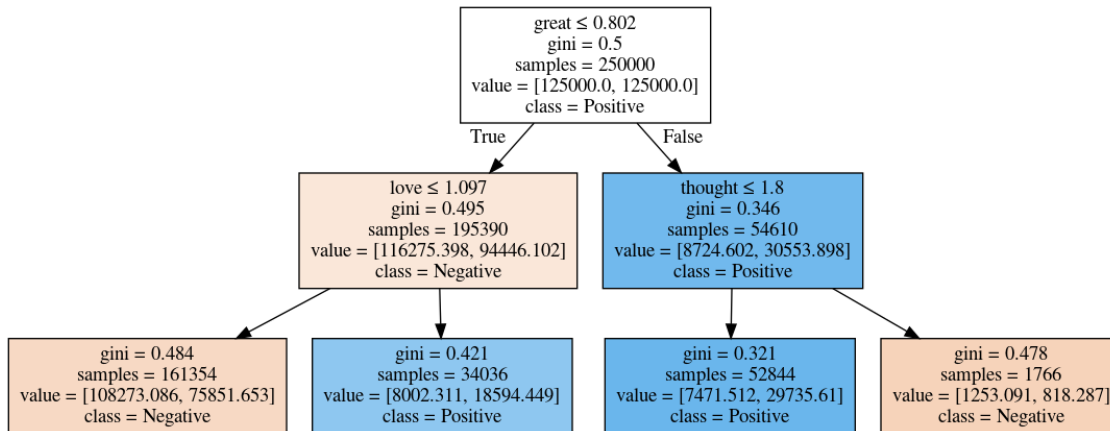
```
[46]: depth = [1, 5, 10, 50, 100, 500, 100]
min_samples_spl = [5, 10, 100, 500]
fig1 = plt.figure(figsize=[12,12])
ax1 = fig1.add_subplot(111,aspect = 'equal')
ax1.add_patch(
    patches.Arrow(0.45,0.5,-0.25,0.25,width=0.3,color='green',alpha = 0.5)
)
ax1.add_patch(
    patches.Arrow(0.5,0.45,0.25,-0.25,width=0.3,color='red',alpha = 0.5)
)

mean_fpr = np.linspace(0,1,100)
for i in depth:
    for j in min_samples_spl:
        classifier = DecisionTreeClassifier(max_depth=i, min_samples_split=j,
        ↪class_weight='balanced')
        model = CalibratedClassifierCV(classifier,cv=5,method = 'isotonic')
        model.fit(X_train_tfidf,y_train)
        mod_probs = model.predict_proba(X_test_tfidf)[: ,1]
        fpr, tpr, thresholds = metrics.roc_curve(y_test, mod_probs)
        auc = metrics.roc_auc_score(y_test, mod_probs)
        plt.plot(fpr,tpr,label="auc_score/max_depth/min_samples_split_
        ↪="+str(auc) + "/" +str(i) + "/" +str(j))
        plt.legend(loc=4)
        plt.plot([0,1],[0,1],linestyle = '--',lw = 2,color = 'black')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC')
plt.legend(loc="lower right")
plt.text(0.32,0.7,'More accurate area',fontsize = 12)
plt.text(0.63,0.4,'Less accurate area',fontsize = 12)
plt.show()
```




```
[47]: classifier = DecisionTreeClassifier(max_depth=2,class_weight='balanced')
classifier.fit(X_train_tfidf,y_train)
feature_name = count_vect.get_feature_names()
target = ['Negative','Positive']
from sklearn.tree import export_graphviz
graph = tree.export_graphviz(classifier,out_file=None, class_names=target,
    ↪,feature_names = feature_name, filled = True,special_characters=True)
# Draw graph
graph = pydotplus.graph_from_dot_data(graph)
# Show graph
Image(graph.create_png())
```


[47]:



```
[48]: classifier = DecisionTreeClassifier(max_depth=15, class_weight='balanced')
classifier.fit(X_train_tfidf,y_train)
#coef = classifier.coef_
probs = classifier.predict_proba(X_test_tfidf)[: ,1]
model = CalibratedClassifierCV(classifier,cv=5,method = 'isotonic')
model.fit(X_train_bow,y_train)
mod_probs = model.predict_proba(X_test_tfidf)[: ,1]
#reliability diagram
fop, mpv = calibration_curve(y_test, probs, n_bins=10,normalize=True)
fop1, mpv1 = calibration_curve(y_test, mod_probs, n_bins=10,normalize=True)
# plot perfectly calibrated
plt.plot([0, 1], [0, 1], linestyle='--',label='perfectly calibrated')
# plot model reliability
plt.plot(mpv, fop, marker='.',label='model uncalibrated')
plt.plot(mpv1, fop1, marker='.',label='model calibrated')
plt.title("Reliability Diagram")
plt.xlabel("mean_predicted_values")
plt.ylabel("fraction_of_positives")
plt.legend()
plt.show()
```



```
[49]: classifier = DecisionTreeClassifier(max_depth=100, min_samples_split=500,
    ↪class_weight='balanced')
classifier.fit(X_train_tfidf,y_train)
model = CalibratedClassifierCV(classifier,cv=5,method = 'isotonic')
model.fit(X_train_tfidf,y_train)
y_pred = model.predict(X_test_tfidf)
acc = f1_score(y_pred,y_test,average='micro')*100
df = pd.DataFrame(confusion_matrix(y_test,y_pred))
sns.heatmap(df,annot=True,fmt="d")
plt.show()
print('\n\nThe accuracy of the DecisionTreeClassifier for depth = %d and
    ↪min_samples_split = %d is %f%%' % (100,500, acc))
```



The accuracy of the DecisionTreeClassifier for depth = 100 and min_samples_split = 500 is 85.543333%

```
[50]: y_pred = model.predict(X_test_tfidf)
      from sklearn.metrics import classification_report
      print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.79	0.25	0.39	5346
1	0.86	0.99	0.92	24654
micro avg	0.86	0.86	0.86	30000
macro avg	0.83	0.62	0.65	30000
weighted avg	0.85	0.86	0.82	30000

```
[56]: importance = classifier.feature_importances_
      class_labels = model.classes_
      feature_names = tfidf_vec.get_feature_names()
      topn_class1 = sorted(zip(importance, feature_names), reverse=False)[:20]
      topn_class2 = sorted(zip(importance, feature_names), reverse=True)[:20]
      print("Important words in negative reviews")
      for importanc, feat in topn_class1:
          print(class_labels[0], importanc, feat)
```

```

print("-----")
print("Important words in positive reviews")
for importanc, feat in topn_class2:
    print(class_labels[1], importanc, feat)

```

Important words in negative reviews

```

0 0.0 aa
0 0.0 aaa
0 0.0 aaaa
0 0.0 aaaaa
0 0.0 aaaaaaaaaa
0 0.0 aaaaaaaaaaaaaa
0 0.0 aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaabr
0 0.0 aaaaaaaaaaaaaaaaaaaaaaargh
0 0.0 aaaaaaaaagghh
0 0.0 aaaaaaahhhhhh
0 0.0 aaaaaarrrrrrggghhh
0 0.0 aaaaaabr
0 0.0 aaaaaah
0 0.0 aaaaaahhh
0 0.0 aaaaaahhhhbr
0 0.0 aaaaaahhhhhyaaaaaa
0 0.0 aaaaaawwwwwwww
0 0.0 aaaaaah
0 0.0 aaaaahhhhhhhhhhhhhhhthe
0 0.0 aaaaawill

```

Important words in positive reviews

```

1 0.10418003844166326 great
1 0.055505336973219896 love
1 0.05124971894124606 best
1 0.03777943866638621 delicious
1 0.026288969232898297 disappointed
1 0.023825810017042738 perfect
1 0.023316608928900848 good
1 0.02036001477438885 favorite
1 0.02014687423002695 bad
1 0.018050280727092333 thought
1 0.017882415172827775 excellent
1 0.013602705438262717 wonderful
1 0.01174013265024442 highly
1 0.011696928089392907 money
1 0.01162862553202851 easy
1 0.01096842614565546 nice
1 0.010109408694509759 awful
1 0.009656989254082569 worst
1 0.0077253510765387156 taste

```

1 0.007634554669081282 terrible

Word2Vec as vectorizer

```
[63]: list_of_sent = []

for sent in data_without_dup['final_string'].values:
    list_of_sent.append(sent.split())
```

```
[64]: from gensim.models import Word2Vec
from gensim.models import KeyedVectors
```

```
[66]: mod = KeyedVectors.load_word2vec_format("/home/niranjan/Downloads/
↳GoogleNews-vectors-negative300.bin", binary=True)
```

```
[101]: X_train_avg_w2v = []
w2v_model = Word2Vec(list_of_sent[0:100000],min_count=5,size=100,workers=4)
w2v_words = list(w2v_model.wv.vocab)

for sent in list_of_sent[0:100000]:
    sent_vec = np.zeros(100)
    count_words = 0
    for words in sent:
        if words in w2v_words:
            vec = w2v_model.wv[words]
            sent_vec += vec
            count_words +=1
    if count_words !=0:
        sent_vec = sent_vec/count_words
    X_train_avg_w2v.append(sent_vec)
```

```
[102]: X_test_avg_w2v = []
w2v_model = Word2Vec(list_of_sent[100000:120000],min_count=5,size=
↳100,workers=4)
w2v_words = list(w2v_model.wv.vocab)
for sent in list_of_sent[100000:120000]:
    sent_vec = np.zeros(100)
    count_words = 0
    for words in sent:
        if words in w2v_words:
            vec = w2v_model.wv[words]
            sent_vec += vec
            count_words +=1
    if count_words !=0:
        sent_vec = sent_vec/count_words
    X_test_avg_w2v.append(sent_vec)
```

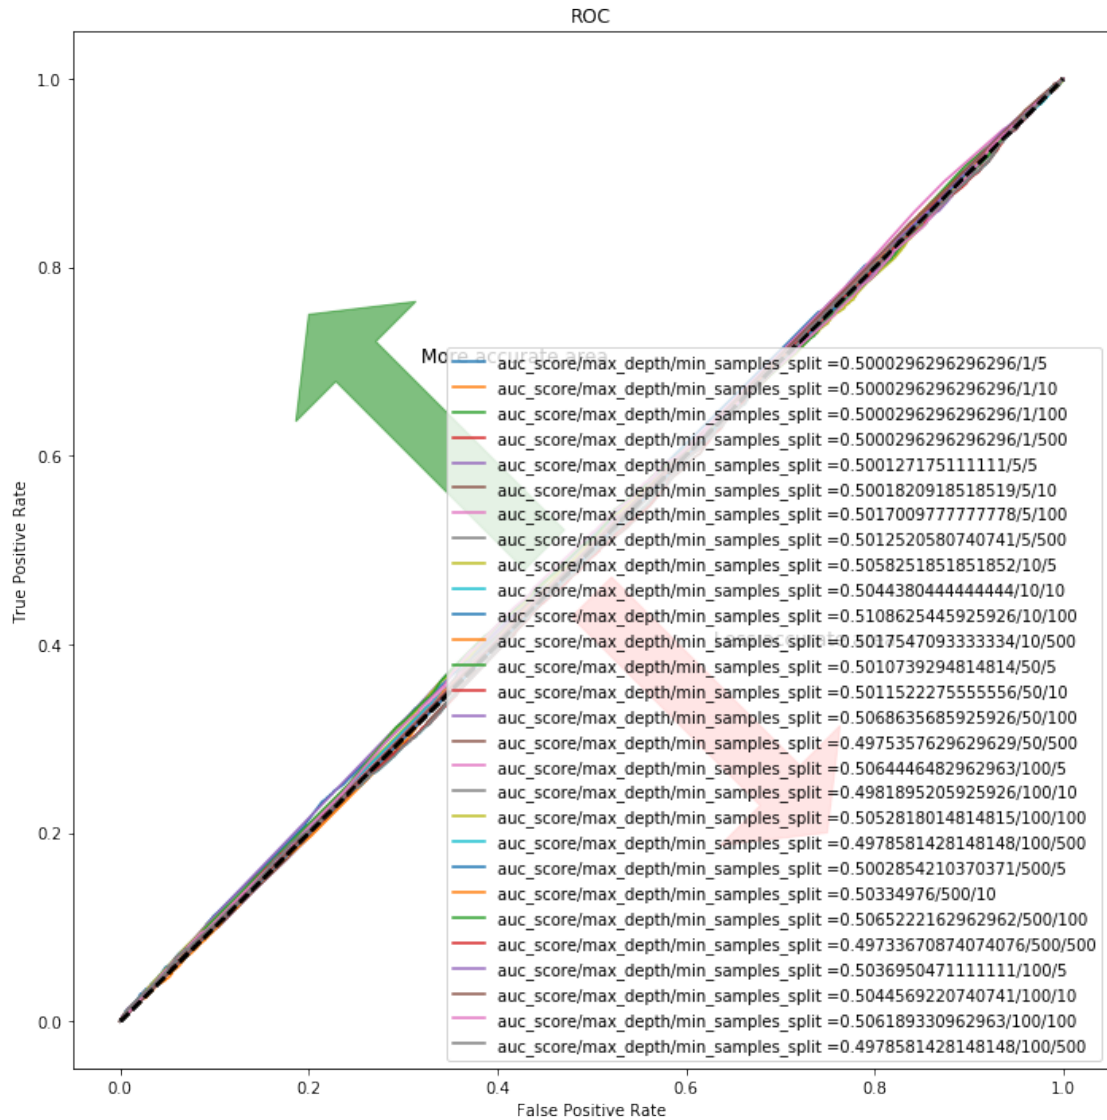


```
[28]: import pickle
pickle_out1 = open("/home/niranjana/Downloads/UBUNTU 18_1/AppliedAI/
↳X_train_avg_w2v","rb")
pickle_out2 = open("/home/niranjana/Downloads/UBUNTU 18_1/AppliedAI/
↳X_cv_avg_w2v","rb")
X_train_avg_w2v = pickle.load(pickle_out1)
X_test_avg_w2v = pickle.load(pickle_out2)
pickle_out1.close()
pickle_out2.close()

[19]: y_train = y[0:100000]
y_test= y[100000:120000]

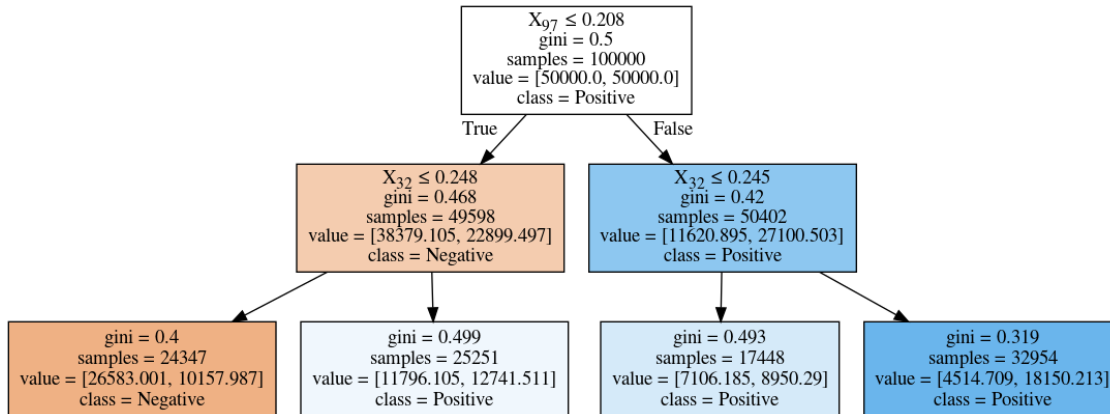
[59]: depth = [1, 5, 10, 50, 100, 500, 100]
min_samples_spl = [5, 10, 100, 500]
fig1 = plt.figure(figsize=[12,12])
ax1 = fig1.add_subplot(111,aspect = 'equal')
ax1.add_patch(
    patches.Arrow(0.45,0.5,-0.25,0.25,width=0.3,color='green',alpha = 0.5)
)
ax1.add_patch(
    patches.Arrow(0.5,0.45,0.25,-0.25,width=0.3,color='red',alpha = 0.5)
)

mean_fpr = np.linspace(0,1,100)
for i in depth:
    for j in min_samples_spl:
        classifier = DecisionTreeClassifier(max_depth=i, min_samples_split=j,
↳class_weight='balanced')
        model = CalibratedClassifierCV(classifier,cv=5,method = 'isotonic')
        model.fit(X_train_avg_w2v,y_train)
        mod_probs = model.predict_proba(X_test_avg_w2v)[:,-1]
        fpr, tpr, thresholds = metrics.roc_curve(y_test, mod_probs)
        auc = metrics.roc_auc_score(y_test, mod_probs)
        plt.plot(fpr,tpr,label="auc_score/max_depth/min_samples_split_
↳="+str(auc) + "/" +str(i) + "/" +str(j))
        plt.legend(loc=4)
        plt.plot([0,1],[0,1],linestyle = '--',lw = 2,color = 'black')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC')
plt.legend(loc="lower right")
plt.text(0.32,0.7,'More accurate area',fontsize = 12)
plt.text(0.63,0.4,'Less accurate area',fontsize = 12)
plt.show()
```




```
[104]: classifier = DecisionTreeClassifier(max_depth=2,class_weight='balanced')
classifier.fit(X_train_avg_w2v,y_train)
#feature_name = count_vect.get_feature_names()
target = ['Negative','Positive']
from sklearn.tree import export_graphviz
graph = tree.export_graphviz(classifier,out_file=None, class_names=target ,
    ↳filled = True,special_characters=True)
# Draw graph
graph = pydotplus.graph_from_dot_data(graph)
# Show graph
Image(graph.create_png())
```

[104]:



```
[29]: classifier = DecisionTreeClassifier(max_depth=15, class_weight='balanced')
classifier.fit(X_train_avg_w2v,y_train)
#coef = classifier.coef_
probs = classifier.predict_proba(X_test_avg_w2v)[: ,1]
model = CalibratedClassifierCV(classifier,cv=5,method = 'isotonic')
model.fit(X_train_avg_w2v,y_train)
mod_probs = model.predict_proba(X_test_avg_w2v)[: ,1]
#reliability diagram
fop, mpv = calibration_curve(y_test, probs, n_bins=10,normalize=True)
fop1, mpv1 = calibration_curve(y_test, mod_probs, n_bins=10,normalize=True)
# plot perfectly calibrated
plt.plot([0, 1], [0, 1], linestyle='--',label='perfectly calibrated')
# plot model reliability
plt.plot(mpv, fop, marker='.',label='model uncaliberated')
plt.plot(mpv1, fop1, marker='.',label='model calibrated')
plt.title("Reliability Diagram")
plt.xlabel("mean_predicted_values")
plt.ylabel("fraction_of_positives")
plt.legend()
plt.show()
```



```
[30]: y_pred = model.predict(X_test_avg_w2v)
      from sklearn.metrics import classification_report
      print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.00	0.00	0.00	3125
1	0.84	1.00	0.92	16875
micro avg	0.84	0.84	0.84	20000
macro avg	0.42	0.50	0.46	20000
weighted avg	0.71	0.84	0.77	20000

tfidf-Word2Vec

```
[105]: # TF-IDF weighted Word2Vec
      vec = TfidfVectorizer()
      vec.fit_transform(X_train[0:100000], y_train[0:100000])
      tfidf_feat = vec.get_feature_names() # tfidf words/col-names
      dictionary = dict(zip(vec.get_feature_names(), list(vec.idf_)))
      # final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val =
      ↪ tfidf
```

```

tfidf_train_vectors = []; # the tfidf-w2v for each sentence/review is stored in
    ↳ this list
row=0;
w2v_model = Word2Vec(list_of_sent[0:100000],min_count=5,size=100,workers=4)
w2v_words = list(w2v_model.wv.vocab)
for sent in tqdm(list_of_sent[0:100000]): # for each review/sentence
    sent_vec = np.zeros(100) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            if word in tfidf_feat:
                vect = w2v_model.wv[word]
#                tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
                # to reduce the computation we are
                # dictionary[word] = idf value of word in whole corpus
                # sent.count(word) = tf value of word in this review
                tf_idf = dictionary[word]*(sent.count(word)/len(sent))
                sent_vec += (vect * tf_idf)
                weight_sum += tf_idf
            else:
                break
    if weight_sum != 0:
        sent_vec /= weight_sum
    tfidf_train_vectors.append(sent_vec)
    row += 1

```

100%| | 100000/100000 [2:06:45<00:00, 8.29it/s]

```

[107]: # TF-IDF weighted Word2Vec
vec = TfidfVectorizer()
vec.fit_transform(X[100000:120000],y[100000:120000])
tfidf_feat = vec.get_feature_names() # tfidf words/col-names
dictionary = dict(zip(vec.get_feature_names(), list(vec.idf_)))
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val =
    ↳ tfidf
tfidf_test_vectors = []; # the tfidf-w2v for each sentence/review is stored in
    ↳ this list
row=0;
w2v_model = Word2Vec(list_of_sent[100000:120000],min_count=5,size=100,workers=4)
w2v_words = list(w2v_model.wv.vocab)
for sent in tqdm(list_of_sent[100000:120000]): # for each review/sentence
    sent_vec = np.zeros(100) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            if word in tfidf_feat:
                vect = w2v_model.wv[word]

```

```
#         tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
# to reduce the computation we are
# dictionary[word] = idf value of word in whole corpus
# sent.count(word) = tf value of word in this review
        tf_idf = dictionary[word]*(sent.count(word)/len(sent))
        sent_vec += (vect * tf_idf)
        weight_sum += tf_idf
    else:
        break
if weight_sum != 0:
    sent_vec /= weight_sum
tfidf_test_vectors.append(sent_vec)
row += 1
```

100%| | 20000/20000 [07:21<00:00, 45.26it/s]

```
[4]: import pickle
pickle_out1 = open("/home/niranjana/Downloads/UBUNTU 18_1/AppliedAI/
↳tfidf_train_vectors", "rb")
pickle_out2 = open("/home/niranjana/Downloads/UBUNTU 18_1/AppliedAI/
↳tfidf_test_vectors", "rb")
tfidf_train_vectors = pickle.load(pickle_out1)
tfidf_test_vectors = pickle.load(pickle_out2)
pickle_out1.close()
pickle_out2.close()
```


```
[84]: depth = [1, 5, 10, 50, 100, 500, 1000]
min_samples_split = [5, 10, 100, 500]
fig1 = plt.figure(figsize=[12,12])
ax1 = fig1.add_subplot(111, aspect = 'equal')
ax1.add_patch(
    patches.Arrow(0.45,0.5,-0.25,0.25,width=0.3,color='green',alpha = 0.5)
)
ax1.add_patch(
    patches.Arrow(0.5,0.45,0.25,-0.25,width=0.3,color='red',alpha = 0.5)
)

mean_fpr = np.linspace(0,1,100)
for i in depth:
    for j in min_samples_split:
        classifier = DecisionTreeClassifier(max_depth=i, min_samples_split=j,
↳class_weight='balanced')
        model = CalibratedClassifierCV(classifier,cv=5,method = 'isotonic')
        model.fit(tfidf_train_vectors,y_train)
        mod_probs = model.predict_proba(tfidf_test_vectors)[:,-1]
        fpr, tpr, thresholds = metrics.roc_curve(y_test, mod_probs)
        auc = metrics.roc_auc_score(y_test, mod_probs)
```

```

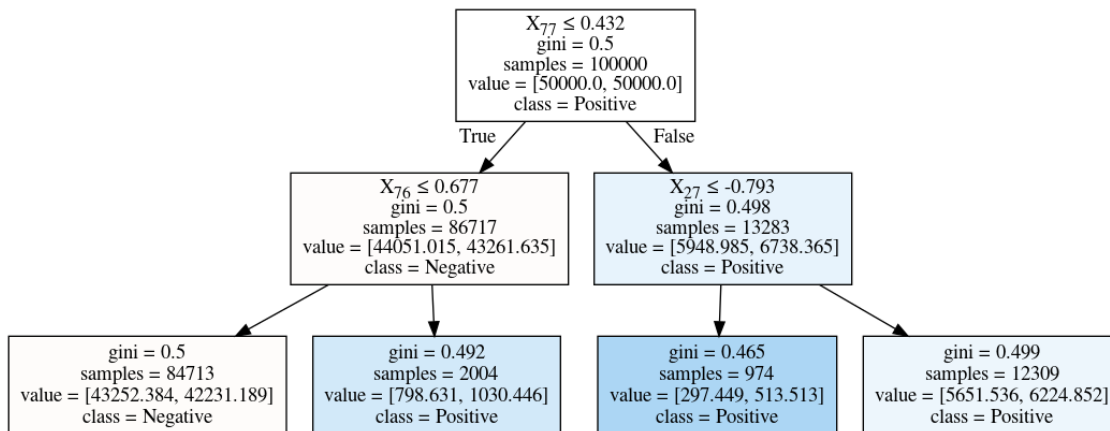
plt.plot(fpr, tpr, label="auc_score/max_depth/min_samples_split_
↪ "+str(auc) + "/" + str(i) + "/" + str(j))
plt.legend(loc=4)
plt.plot([0,1],[0,1], linestyle = '--', lw = 2, color = 'black')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC')
plt.legend(loc="lower right")
plt.text(0.32, 0.7, 'More accurate area', fontsize = 12)
plt.text(0.63, 0.4, 'Less accurate area', fontsize = 12)
plt.show()

```




```
[24]: classifier = DecisionTreeClassifier(max_depth=2,class_weight='balanced')
classifier.fit(tfidf_train_vectors,y_train)
target = ['Negative','Positive']
from sklearn.tree import export_graphviz
graph = tree.export_graphviz(classifier,out_file=None, class_names=target,
    ↳filled = True,special_characters=True)
# Draw graph
graph = pydotplus.graph_from_dot_data(graph)
# Show graph
Image(graph.create_png())
```


[24]:



```
[94]: classifier = DecisionTreeClassifier(max_depth=15, class_weight='balanced')
classifier.fit(tfidf_train_vectors,y_train)
#coef = classifier.coef_
probs = classifier.predict_proba(X_test_avg_w2v)[: ,1]
model = CalibratedClassifierCV(classifier,cv=5,method = 'isotonic')
model.fit(tfidf_train_vectors,y_train)
mod_probs = model.predict_proba(tfidf_test_vectors)[: ,1]
#reliability diagram
fop, mpv = calibration_curve(y_test, probs, n_bins=10,normalize=True)
fop1, mpv1 = calibration_curve(y_test, mod_probs, n_bins=10,normalize=True)
# plot perfectly calibrated
plt.plot([0, 1], [0, 1], linestyle='--',label='perfectly calibrated')
# plot model reliability
plt.plot(mpv, fop, marker='.',label='model uncalibrated')
plt.plot(mpv1, fop1, marker='.',label='model calibrated')
plt.title("Reliability Diagram")
plt.xlabel("mean_predicted_values")
plt.ylabel("fraction_of_positives")
plt.legend()
plt.show()
```

```
[124]: classifier = DecisionTreeClassifier(max_depth=100, min_samples_split=500,
    ↳class_weight='balanced')
classifier.fit(tfidf_train_vectors,y_train)
model = CalibratedClassifierCV(classifier,cv=5,method = 'isotonic')
model.fit(tfidf_train_vectors,y_train)
y_pred = model.predict(tfidf_test_vectors)
acc = f1_score(y_pred,y_test,average='micro')*100
df = pd.DataFrame(confusion_matrix(y_test,y_pred))
sns.heatmap(df,annot=True,fmt="d")
plt.show()
print('\n\nThe accuracy of the DecisionTreeClassifier for depth = %d and
    ↳min_samples_split = %d is %f%%' % (100,500, acc))
```



The accuracy of the DecisionTreeClassifier for depth = 100 and min_samples_split = 500 is 84.375000%

```
[125]: y_pred = model.predict(tfidf_test_vectors)
from sklearn.metrics import classification_report
print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
0	0.00	0.00	0.00	3125
1	0.84	1.00	0.92	16875
micro avg	0.84	0.84	0.84	20000
macro avg	0.42	0.50	0.46	20000
weighted avg	0.71	0.84	0.77	20000

```
/home/niranjan/anaconda3/lib/python3.6/site-
packages/sklearn/metrics/classification.py:1143: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples.
```

```
'precision', 'predicted', average, warn_for)
/home/niranjan/anaconda3/lib/python3.6/site-
packages/sklearn/metrics/classification.py:1143: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
```

predicted samples.

```
'precision', 'predicted', average, warn_for)
/home/niranjana/anaconda3/lib/python3.6/site-
packages/sklearn/metrics/classification.py:1143: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels with no
predicted samples.
'precision', 'predicted', average, warn_for)
```

```
[19]: from prettytable import PrettyTable
x = PrettyTable()
x.add_column("important_features_class_[0]",topn_class1)
x.add_column("important_features_class_[1]",topn_class2)
y = PrettyTable()
y.
  ↳field_names=["Vectorizer","Model","depth","min_samples_split","Auc_score","f1_score(microq
  ↳average)"]
y.add_row(["BOW","DecisionTreeClassifier","100","500",".8668","0.86"])
y.add_row(["tfidf","DecisionTreeClassifier","100","500",".8666","0.86"])
y.add_row(["Word2Vec","DecisionTreeClassifier","10","100",".5108","0.84"])
y.add_row(["tfidf-Word2Vec","DecisionTreeClassifier","5","500",".5001","0.84"])
z = PrettyTable()
z.field_names = ["Vectorizer","Model","Precision", "Recall"]
z.add_row(["BOW","DecisionTreeClassifier",0.86,0.86])
z.add_row(["tfidf","DecisionTreeClassifier",0.86,0.86])
z.add_row(["Word2Vec","DecisionTreeClassifier",0.84,0.84])
z.add_row(["tfidf-Word2Vec","DecisionTreeClassifier",0.84,0.84])

print(x)
print(y)
print(z)
```

```
+-----+-----+
-----+
|           important_features_class_[0]           |
important_features_class_[1] |
+-----+-----+
-----+
|           (0.0, 'aa')           |
(0.0990350154071816, 'great') |
|           (0.0, 'aaa')           |
(0.05358931950541775, 'love') |
|           (0.0, 'aaaa')           |
(0.04721148244145879, 'best') |
|           (0.0, 'aaaaa')           |
(0.036247023890801, 'delicious') |
|           (0.0, 'aaaaaaaaa')           |
(0.03024856779790211, 'disappointed') |
|           (0.0, 'aaaaaaaaaaaaaaaaa')           |
```

```

(0.026546502488069848, 'bad') |
| (0.0, 'aaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaaabr') |
(0.022347812312398498, 'perfect') |
| (0.0, 'aaaaaaaaaaaaaaaaaaaaargh') |
(0.020542713887296185, 'thought') |
| (0.0, 'aaaaaaaagghh') |
(0.019491600028405522, 'good') |
| (0.0, 'aaaaaaahhhhhh') |
(0.019195535082647958, 'favorite') |
| (0.0, 'aaaaaaarrrrrggghhh') |
(0.018327058204032076, 'excellent') |
| (0.0, 'aaaaabr') |
(0.016600775553961878, 'would') |
| (0.0, 'aaaaah') |
(0.014625911816804999, 'money') |
| (0.0, 'aaaaahhh') |
(0.01213348699346732, 'worst') |
| (0.0, 'aaaaahhhhbr') |
(0.011647719525502632, 'wonderful') |
| (0.0, 'aaaaahhhhhhyaaaaaa') |
(0.011115672121797858, 'awful') |
| (0.0, 'aaaaaawwwwwwwww') |
(0.010046053897243981, 'highly') |
| (0.0, 'aaaaah') |
(0.009916758070424038, 'review') |
| (0.0, 'aaaaahhhhhhhhhhhhhhhthe') |
(0.009036311856948726, 'unfortunately') |
| (0.0, 'aaaaawill') |
(0.00901390900928107, 'nice') |
+-----+
-----+
+-----+-----+-----+-----+
--+-----+
| Vectorizer | Model | depth | min_samples_split |
Auc_score | f1_score(micro average) |
+-----+-----+-----+-----+
--+-----+
| BOW | DecisionTreeClassifier | 100 | 500 | .8668
| 0.86 |
| tfidf | DecisionTreeClassifier | 100 | 500 | .8666
| 0.86 |
| Word2Vec | DecisionTreeClassifier | 10 | 100 | .5108
| 0.84 |
| tfidf-Word2Vec | DecisionTreeClassifier | 5 | 500 | .5001
| 0.84 |
+-----+-----+-----+-----+
--+-----+
+-----+-----+-----+-----+

```

Vectorizer	Model	Precision	Recall
BOW	DecisionTreeClassifier	0.86	0.86
tfidf	DecisionTreeClassifier	0.86	0.86
Word2Vec	DecisionTreeClassifier	0.84	0.84
tfidf-Word2Vec	DecisionTreeClassifier	0.84	0.84