

# KnnModelForAmazonFoodReview (1)

April 7, 2020

KNN model implementation for amazon review dataset using BOW, tf-idf, Avg Word2Vec, tfidf- Word2Vec

```
[1]: import numpy as np
import pandas as pd
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.feature_extraction.text import TfidfVectorizer
import sklearn
from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import train_test_split
import sqlite3
from tqdm import tqdm
import warnings
warnings.filterwarnings('ignore')
import string
from sklearn.model_selection import TimeSeriesSplit
from sklearn.model_selection import cross_val_score
from sklearn import cross_validation
from sklearn.metrics import accuracy_score
from sklearn.decomposition import TruncatedSVD
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn import metrics
from sklearn.metrics import confusion_matrix
from sklearn.model_selection import GridSearchCV
from sklearn.metrics import f1_score
import pickle
%matplotlib inline

[2]: con = sqlite3.connect('/home/niranjan/Downloads/database.sqlite')
data = pd.read_sql_query('select * from Reviews where Score!=3 ',con)
data['Score'] = [1 if i > 3 else 0 for i in data['Score']]
#print(data.shape)

[3]: data_sort = data.sort_values(by='ProductId',ascending=True, inplace=False,
    ↪kind='quicksort')
data_groupby = data_sort[['ProductId', 'Score']].groupby('Score').count()
```

```
#print(data_groupby)
```

## Removing duplicates data

```
[4]: data_without_dup = data_sort.  
      ↳ drop_duplicates(subset={'UserId', 'ProfileName', 'Time', 'Text'}, keep='first', inplace=False)  
      #print(data_without_dup.shape)  
      data_groupby = data_without_dup[['ProductId', 'Score']].groupby('Score').count()  
      #print(data_groupby)
```

```
[5]: data_without_dup = data_without_dup[data_without_dup.  
      ↳ HelpfulnessNumerator<=data_without_dup.HelpfulnessDenominator]  
      #print(data_without_dup.shape)
```

```
[6]: import nltk  
      nltk.download('stopwords')  
      nltk.download('wordnet')  
      from nltk.corpus import stopwords  
      from nltk import WordNetLemmatizer  
      stop = set(stopwords.words('english'))  
      lemma = nltk.WordNetLemmatizer()
```

```
[nltk_data] Downloading package stopwords to  
[nltk_data]      /home/niranjana/nltk_data...  
[nltk_data] Package stopwords is already up-to-date!  
[nltk_data] Downloading package wordnet to /home/niranjana/nltk_data...  
[nltk_data] Package wordnet is already up-to-date!
```

```
[7]: import re  
      def cleanhtml(words):  
          tag = re.compile(r'<.*?>')  
          cleanSent = re.sub(tag, '', words)  
          return cleanSent  
  
      def PuncRemov(words):  
          tag = re.compile(r'[^a-zA-Z]')  
          cleanSent = re.sub(tag, '', words)  
          return cleanSent
```

\*\*\*\*\_data cleaning\_\*\*\*\*

---

removal of html tags, symbols other than alphabets, stopwords and performing lemmatization as part of data pre-processing.

---

**Lemmatization :-** is the process of grouping together the inflected forms of a word so they can be analysed as a single item, identified by the word's lemma, or dictionary

form.\*\*

```
[8]: final_string = []
    str1 = ' '
    positive_word = []
    negative_word = []
    i=0
    for sen in data_without_dup['Text'].values:
        filtered_word = []
        sent = cleanhtml(sen)
        for word in sent.split():
            cleanWord = PuncRemov(word)
            for cleaned_words in cleanWord.split():
                if ((len(cleaned_words) > 2) & (cleaned_words.isalpha())):
                    if (cleaned_words.islower() not in stop):
                        w = (lemma.lemmatize(cleaned_words.lower())).encode('utf8')
                        filtered_word.append(w)
                        if data_without_dup['Score'].values[i]=='positive':
                            positive_word.append(w)
                        else:
                            negative_word.append(w)
                    else:
                        continue
                else:
                    continue
            str1 = b" ".join(filtered_word)
            final_string.append(str1)
            i = i+1
```

```
[9]: data_without_dup['cleaned_text'] = final_string
    data_without_dup['cleaned_text'] = data_without_dup['cleaned_text'].str.
        ↳ decode('utf8')
```

sorting data based on time in ascending order

```
[10]: data_without_dup = data_without_dup.
        ↳ sort_values(by='Time', ascending=True, inplace=False)
```

```
[11]: X = data_without_dup['cleaned_text']
    y = data_without_dup['Score']
```

```
[12]: X_train= X[0:100000]
    y_train = y[0:100000]
    X_cv= X[100000:120000]
    y_cv = y[100000:120000]
    X_test= X[120000:140000]
    y_test = y[120000:140000]
```

```
[13]: print("null_accuracy :",max(y.mean(),(1-y.mean())))
```

null\_accuracy : 0.843178067446337

pre-processing of text using BOW

```
[21]: count_vect = CountVectorizer()
X_train_bow = count_vect.fit_transform(X_train)
X_test_bow = count_vect.transform(X_test)
X_cv_bow = count_vect.transform(raw_documents=X_cv)
print(X_cv_bow.shape)
```

(20000, 88686)

```
[ ]: import pickle
pickle_in1 = open("/home/niranjan/Downloads/AppliedAI/X_train_bow","rb")
pickle_in2 = open("/home/niranjan/Downloads/AppliedAI/X_cv_bow","rb")
pickle_in3 = open("/home/niranjan/Downloads/AppliedAI/X_test_bow","rb")

X_train_bow = pickle.load(pickle_in1)
X_cv_bow = pickle.load(pickle_in2)
X_test_bow = pickle.load(pickle_in3)
pickle_in1.close()
pickle_in2.close()
pickle_in3.close()
```

Best fit using Cross validation

```
[74]: myList = list(range(0,50))

#creating odd list for k in Knn
neighbor = [x for x in myList if x%2 != 0]

cv_scores = []
for k in neighbor:
    classifier = KNeighborsClassifier(n_neighbors=k,algorithm='kd_tree',n_jobs=1)
    classifier.fit(X_train_bow,y_train)
    for i in range(0,X_cv_bow.shape[0],1000):
        y_pred = []
        y_pred.append(classifier.predict(X_cv_bow[i:i+1000]))
        y_pred_flat = [item for sublist in y_pred for item in sublist]
        cv_scores.append(f1_score(y_cv[i:i+1000],y_pred_flat,average='micro'))

#Determining optimal value of k

optimal_k = neighbor[cv_scores.index(max(cv_scores))]
print("optimal number of is: ", optimal_k)
```

```
plt.plot(neighbor,cv_scores)
for xy in zip(neighbor, np.round(cv_scores,3)):
    plt.annotate('%s, %s)' % xy, xy=xy, textcoords='data')

plt.xlabel('Number of Neighbors K')
plt.ylabel('f1_score')
plt.show()

# print("the misclassification error for each k value is : ", np.round(MSE,3))
```

```
/home/niranjana/anaconda3/lib/python3.7/site-
packages/sklearn/neighbors/base.py:212: UserWarning: cannot use tree with sparse
input: using brute force
    warnings.warn("cannot use tree with sparse input: "
/home/niranjana/anaconda3/lib/python3.7/site-
packages/sklearn/neighbors/base.py:212: UserWarning: cannot use tree with sparse
input: using brute force
    warnings.warn("cannot use tree with sparse input: "
/home/niranjana/anaconda3/lib/python3.7/site-
packages/sklearn/neighbors/base.py:212: UserWarning: cannot use tree with sparse
input: using brute force
    warnings.warn("cannot use tree with sparse input: "
/home/niranjana/anaconda3/lib/python3.7/site-
packages/sklearn/neighbors/base.py:212: UserWarning: cannot use tree with sparse
input: using brute force
    warnings.warn("cannot use tree with sparse input: "
/home/niranjana/anaconda3/lib/python3.7/site-
packages/sklearn/neighbors/base.py:212: UserWarning: cannot use tree with sparse
input: using brute force
    warnings.warn("cannot use tree with sparse input: "
/home/niranjana/anaconda3/lib/python3.7/site-
packages/sklearn/neighbors/base.py:212: UserWarning: cannot use tree with sparse
input: using brute force
    warnings.warn("cannot use tree with sparse input: "
/home/niranjana/anaconda3/lib/python3.7/site-
packages/sklearn/neighbors/base.py:212: UserWarning: cannot use tree with sparse
input: using brute force
    warnings.warn("cannot use tree with sparse input: "
/home/niranjana/anaconda3/lib/python3.7/site-
packages/sklearn/neighbors/base.py:212: UserWarning: cannot use tree with sparse
input: using brute force
    warnings.warn("cannot use tree with sparse input: "
```

```


/home/niranjana/anaconda3/lib/python3.7/site-
packages/sklearn/neighbors/base.py:212: UserWarning: cannot use tree with sparse
input: using brute force
    warnings.warn("cannot use tree with sparse input: "
/home/niranjana/anaconda3/lib/python3.7/site-
packages/sklearn/neighbors/base.py:212: UserWarning: cannot use tree with sparse
input: using brute force
    warnings.warn("cannot use tree with sparse input: "
/home/niranjana/anaconda3/lib/python3.7/site-
packages/sklearn/neighbors/base.py:212: UserWarning: cannot use tree with sparse
input: using brute force
    warnings.warn("cannot use tree with sparse input: "
/home/niranjana/anaconda3/lib/python3.7/site-
packages/sklearn/neighbors/base.py:212: UserWarning: cannot use tree with sparse
input: using brute force
    warnings.warn("cannot use tree with sparse input: "
/home/niranjana/anaconda3/lib/python3.7/site-
packages/sklearn/neighbors/base.py:212: UserWarning: cannot use tree with sparse
input: using brute force
    warnings.warn("cannot use tree with sparse input: "
/home/niranjana/anaconda3/lib/python3.7/site-
packages/sklearn/neighbors/base.py:212: UserWarning: cannot use tree with sparse
input: using brute force
    warnings.warn("cannot use tree with sparse input: "
/home/niranjana/anaconda3/lib/python3.7/site-
packages/sklearn/neighbors/base.py:212: UserWarning: cannot use tree with sparse
input: using brute force
    warnings.warn("cannot use tree with sparse input: "
/home/niranjana/anaconda3/lib/python3.7/site-
packages/sklearn/neighbors/base.py:212: UserWarning: cannot use tree with sparse
input: using brute force
    warnings.warn("cannot use tree with sparse input: "
/home/niranjana/anaconda3/lib/python3.7/site-
packages/sklearn/neighbors/base.py:212: UserWarning: cannot use tree with sparse
input: using brute force
    warnings.warn("cannot use tree with sparse input: "
/home/niranjana/anaconda3/lib/python3.7/site-
packages/sklearn/neighbors/base.py:212: UserWarning: cannot use tree with sparse
input: using brute force
    warnings.warn("cannot use tree with sparse input: "

```

```

/home/niranjana/anaconda3/lib/python3.7/site-
packages/sklearn/neighbors/base.py:212: UserWarning: cannot use tree with sparse
input: using brute force
  warnings.warn("cannot use tree with sparse input: "
/home/niranjana/anaconda3/lib/python3.7/site-
packages/sklearn/neighbors/base.py:212: UserWarning: cannot use tree with sparse
input: using brute force
  warnings.warn("cannot use tree with sparse input: "
/home/niranjana/anaconda3/lib/python3.7/site-
packages/sklearn/neighbors/base.py:212: UserWarning: cannot use tree with sparse
input: using brute force
  warnings.warn("cannot use tree with sparse input: "
/home/niranjana/anaconda3/lib/python3.7/site-
packages/sklearn/neighbors/base.py:212: UserWarning: cannot use tree with sparse
input: using brute force
  warnings.warn("cannot use tree with sparse input: "
optimal number of is: 5

```



```

[153]: y_pred = []
        clf = KNeighborsClassifier(n_neighbors=optimal_k)
        clf.fit(X_train_bow,y_train)
        for i in range(0,X_test_bow.shape[0],1000):
            y_pred.append(clf.predict(X_test_bow[i:i+1000]))
        # evaluate accuracy
        y_pred_flat = [item for sublist in y_pred for item in sublist]

```


```

acc = f1_score(y_test, y_pred_flat,average='micro')*100
print('\nThe accuracy of the knn classifier for k = %d is %f%%' % (optimal_k,
    ↪acc))
fpr, tpr, thresholds = metrics.roc_curve(y_test, y_pred_flat, pos_label=2)
df = pd.DataFrame(confusion_matrix(y_test,y_pred_flat))
sns.heatmap(df,annot=True,fmt="d")
plt.show()

```

The accuracy of the knn classifier for k = 5 is 83.500000%

/home/niranjan/anaconda3/lib/python3.7/site-packages/sklearn/metrics/ranking.py:571: UndefinedMetricWarning: No positive samples in y\_true, true positive value should be meaningless  
UndefinedMetricWarning)



```

[144]: y_pred_proba = []
        for i in range(0,X_test_bow.shape[0],1000):
            y_pred_proba.append(clf.predict_proba(X_test_bow[i:i+1000]))

```

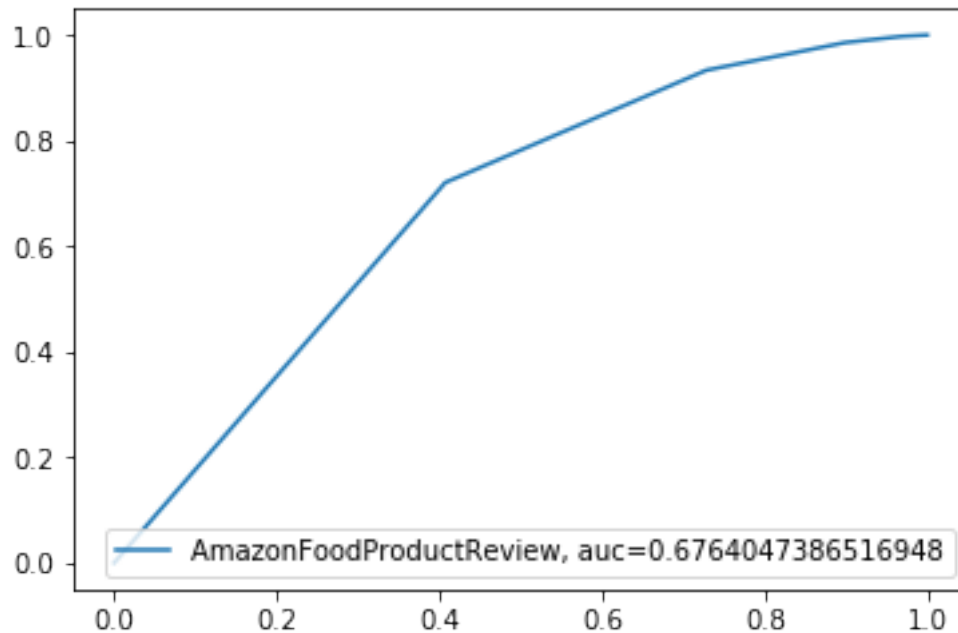
```

[145]: y_pred = []
        for i in y_pred_proba:
            for j in i:
                y_pred.append(j[1])

```



```
[146]: fpr, tpr, thresholds = metrics.roc_curve(y_test, y_pred)
auc = metrics.roc_auc_score(y_test, y_pred)
plt.plot(fpr,tpr,label="AmazonFoodProductReview, auc="+str(auc))
plt.legend(loc=4)
plt.show()
```



```
[147]: myList = list(range(0,50))

#creating odd list for k in Knn
neighbor = [x for x in myList if x%2 != 0]

cv_scores = []
for k in neighbor:
    classifier = KNeighborsClassifier(n_neighbors=k,algorithm='brute',n_jobs=1)
    classifier.fit(X_train_bow,y_train)
    for i in range(0,X_cv_bow.shape[0],1000):
        y_pred = []
        y_pred.append(classifier.predict(X_cv_bow[i:i+1000]))
        y_pred_flat = [item for sublist in y_pred for item in sublist]
        cv_scores.append(f1_score(y_cv[i:i+1000],y_pred_flat,average='micro'))

#Determining optimal value of k

optimal_k_brute = neighbor[cv_scores.index(max(cv_scores))]
```

```

print("optimal number of is: ", optimal_k_brute)

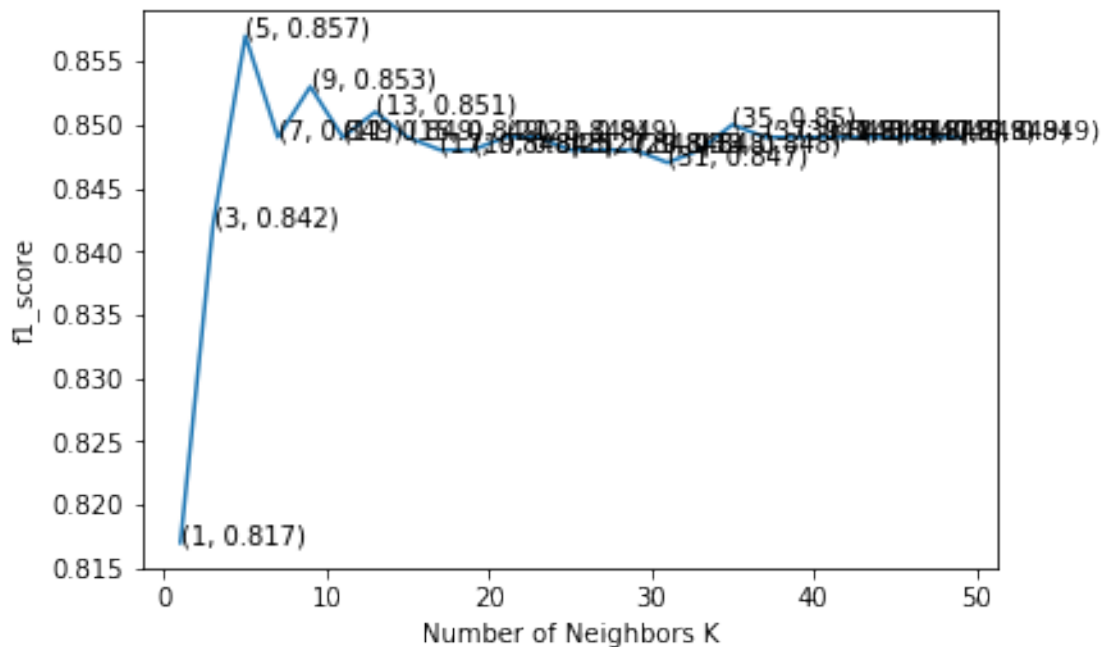
plt.plot(neighbor,cv_scores)
for xy in zip(neighbor, np.round(cv_scores,3)):
    plt.annotate('%s, %s' % xy, xy=xy, textcoords='data')

plt.xlabel('Number of Neighbors K')
plt.ylabel('f1_score')
plt.show()

# print("f1_score for each k value is : ", np.round(cv_scores,3))

```

optimal number of is: 5



```


[152]: y_pred = []
        clf = KNeighborsClassifier(n_neighbors=optimal_k_brute)
        clf.fit(X_train_bow,y_train)
        for i in range(0,X_test_bow.shape[0],1000):
            y_pred.append(clf.predict(X_test_bow[i:i+1000]))
        # evaluate accuracy
        y_pred_flat = [item for sublist in y_pred for item in sublist]
        acc = f1_score(y_test, y_pred_flat,average='micro')*100
        print('\nThe accuracy of the knn classifier for k = %d is %f%%' %
            ↳(optimal_k_brute, acc))
        fpr, tpr, thresholds = metrics.roc_curve(y_test, y_pred_flat, pos_label=2)

```

```
df = pd.DataFrame(confusion_matrix(y_test,y_pred_flat))
sns.heatmap(df,annot=True,fmt="d")
plt.show()
```

The accuracy of the knn classifier for k = 5 is 83.500000%


```
/home/niranjana/anaconda3/lib/python3.7/site-
packages/sklearn/metrics/ranking.py:571: UndefinedMetricWarning: No positive
samples in y_true, true positive value should be meaningless
UndefinedMetricWarning)
```



```
[149]: y_pred_proba = []
for i in range(0,X_test_bow.shape[0],1000):
    y_pred_proba.append(clf.predict_proba(X_test_bow[i:i+1000]))
```

```
[150]: y_pred = []
for i in y_pred_proba:
    for j in i:
        y_pred.append(j[1])
```

```
[151]: fpr, tpr, thresholds = metrics.roc_curve(y_test, y_pred)
auc = metrics.roc_auc_score(y_test, y_pred)
plt.plot(fpr,tpr,label="AmazonFoodProductReview, auc="+str(auc))
plt.legend(loc=4)
plt.show()
```



### Preprocessing using tf-idf

```
[154]: vec = TfidfVectorizer()
```

```
[157]: X_train_tfidf = vec.fit_transform(X_train,y_train)
X_cv_tfidf = vec.transform(X_cv)
X_test_tfidf = vec.transform(X_test)
```

```
[163]: myList = list(range(0,50))

#creating odd list for k in Knn
neighbor = [x for x in myList if x%2 != 0]

cv_scores = []
for k in neighbor:
    classifier = KNeighborsClassifier(n_neighbors=k,algorithm='brute',n_jobs=1)
    classifier.fit(X_train_tfidf,y_train)
    for i in range(0,X_cv_tfidf.shape[0],1000):
        y_pred = []
        y_pred.append(classifier.predict(X_cv_tfidf[i:i+1000]))
        y_pred_flat = [item for sublist in y_pred for item in sublist]
        cv_scores.append(f1_score(y_cv[i:i+1000],y_pred_flat,average='micro'))

#Determining optimal value of k
```

```

optimal_k_brute = neighbor[cv_scores.index(max(cv_scores))]
print("optimal number of is: ", optimal_k_brute)


plt.plot(neighbor,cv_scores)
for xy in zip(neighbor, np.round(cv_scores,3)):
    plt.annotate('%s, %s)' % xy, xy=xy, textcoords='data')

plt.xlabel('Number of Neighbors K')
plt.ylabel('f1_score')
plt.show()

# print("f1_score for each k value is : ", np.round(cv_scores,3))

```

optimal number of is: 15



```


[164]: y_pred = []
clf = KNeighborsClassifier(n_neighbors=optimal_k_brute)
clf.fit(X_train_tfidf,y_train)
for i in range(0,X_test_tfidf.shape[0],1000):
    y_pred.append(clf.predict(X_test_tfidf[i:i+1000]))
# evaluate accuracy
y_pred_flat = [item for sublist in y_pred for item in sublist]
acc = f1_score(y_test, y_pred_flat,average='micro')*100
print('\nThe accuracy of the knn classifier for k = %d is %f%%' %_
↪(optimal_k_brute, acc))

```

```
fpr, tpr, thresholds = metrics.roc_curve(y_test, y_pred_flat, pos_label=2)
df = pd.DataFrame(confusion_matrix(y_test, y_pred_flat))
sns.heatmap(df, annot=True, fmt="d")
plt.show()
```

The accuracy of the knn classifier for k = 15 is 84.360000%

```
/home/niranjana/anaconda3/lib/python3.7/site-
packages/sklearn/metrics/ranking.py:571: UndefinedMetricWarning: No positive
samples in y_true, true positive value should be meaningless
UndefinedMetricWarning)
```




```
[165]: y_pred_proba = []
for i in range(0, X_test_tfidf.shape[0], 1000):
    y_pred_proba.append(clf.predict_proba(X_test_tfidf[i:i+1000]))
```

```
[166]: y_pred = []
for i in y_pred_proba:
    for j in i:
        y_pred.append(j[1])
```

```
[167]: fpr, tpr, thresholds = metrics.roc_curve(y_test, y_pred)
auc = metrics.roc_auc_score(y_test, y_pred)
plt.plot(fpr, tpr, label="AmazonFoodProductReview", auc=str(auc))
plt.legend(loc=4)
```

```
plt.show()
```



### Algorithm as kd\_tree

```
[168]: myList = list(range(0,50))

#creating odd list for k in Knn
neighbor = [x for x in myList if x%2 != 0]

cv_scores = []
for k in neighbor:
    classifier = KNeighborsClassifier(n_neighbors=k,algorithm='kd_tree',n_jobs=1)
    classifier.fit(X_train_tfidf,y_train)
    for i in range(0,X_cv_tfidf.shape[0],1000):
        y_pred = []
        y_pred.append(classifier.predict(X_cv_tfidf[i:i+1000]))
        y_pred_flat = [item for sublist in y_pred for item in sublist]
        cv_scores.append(f1_score(y_cv[i:i+1000],y_pred_flat,average='micro'))

#Determining optimal value of k

optimal_k_kd_tree = neighbor[cv_scores.index(max(cv_scores))]
print("optimal number of k is: ", optimal_k_kd_tree)

plt.plot(neighbor,cv_scores)
```



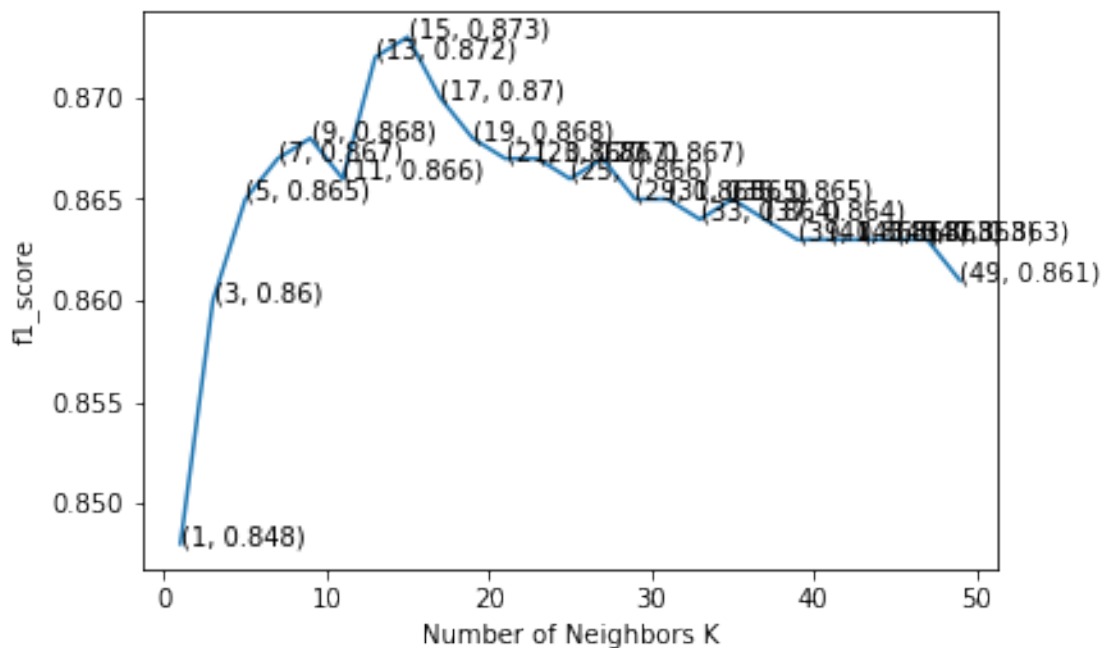




```

input: using brute force
warnings.warn("cannot use tree with sparse input: "
/home/niranjana/anaconda3/lib/python3.7/site-
packages/sklearn/neighbors/base.py:212: UserWarning: cannot use tree with sparse
input: using brute force
warnings.warn("cannot use tree with sparse input: "
/home/niranjana/anaconda3/lib/python3.7/site-
packages/sklearn/neighbors/base.py:212: UserWarning: cannot use tree with sparse
input: using brute force
warnings.warn("cannot use tree with sparse input: "
/home/niranjana/anaconda3/lib/python3.7/site-
packages/sklearn/neighbors/base.py:212: UserWarning: cannot use tree with sparse
input: using brute force
warnings.warn("cannot use tree with sparse input: "
optimal number of is: 15

```



```

[169]: y_pred = []
        clf = KNeighborsClassifier(n_neighbors=optimal_k_kd_tree)
        clf.fit(X_train_tfidf,y_train)
        for i in range(0,X_test_tfidf.shape[0],1000):
            y_pred.append(clf.predict(X_test_tfidf[i:i+1000]))
        # evaluate accuracy
        y_pred_flat = [item for sublist in y_pred for item in sublist]
        acc = f1_score(y_test, y_pred_flat,average='micro')*100

```


```

print('\n\nThe accuracy of the knn classifier for k = %d is %f%%' % (
    ↪(optimal_k_kd_tree, acc))
fpr, tpr, thresholds = metrics.roc_curve(y_test, y_pred_flat, pos_label=2)
df = pd.DataFrame(confusion_matrix(y_test,y_pred_flat))
sns.heatmap(df,annot=True,fmt="d")
plt.show()

```

The accuracy of the knn classifier for k = 15 is 84.360000%

/home/niranjana/anaconda3/lib/python3.7/site-packages/sklearn/metrics/ranking.py:571: UndefinedMetricWarning: No positive samples in y\_true, true positive value should be meaningless  
UndefinedMetricWarning)



```

[170]: y_pred_proba = []
for i in range(0,X_test_tfidf.shape[0],1000):
    y_pred_proba.append(clf.predict_proba(X_test_tfidf[i:i+1000]))

```

```

[171]: y_pred = []
for i in y_pred_proba:
    for j in i:
        y_pred.append(j[1])


```

```

[172]: fpr, tpr, thresholds = metrics.roc_curve(y_test, y_pred)
auc = metrics.roc_auc_score(y_test, y_pred)

```

```
plt.plot(fpr,tpr,label="AmazonFoodProductReview, auc="+str(auc))
plt.legend(loc=4)
plt.show()
```



```
[14]: list_of_sent = []

for sent in data_without_dup['cleaned_text'].values:
    list_of_sent.append(sent.split())
```

```
[15]: from gensim.models import Word2Vec
from gensim.models import KeyedVectors
```

```
[16]: model = KeyedVectors.load_word2vec_format('/home/niranjana/Downloads/data',
↪binary=True)
```

```
[33]: X_train_avg_w2v = []
w2v_model = Word2Vec(list_of_sent[0:100000],min_count=5,size=100,workers=4)
w2v_words = list(w2v_model.wv.vocab)

for sent in list_of_sent[0:100000]:
    sent_vec = np.zeros(100)
    count_words = 0
    for words in sent:
        if words in w2v_words:
            vec = w2v_model.wv[words]
```

```

        sent_vec += vec
        count_words +=1
    if count_words !=0:
        sent_vec = sent_vec/count_words
    X_train_avg_w2v.append(sent_vec)

```

/home/niranjana/anaconda3/lib/python3.7/site-packages/gensim/models/base\_any2vec.py:743: UserWarning: C extension not loaded, training will be slow. Install a C compiler and reinstall gensim for fast training.

"C extension not loaded, training will be slow. "

```

[34]: X_cv_avg_w2v = []
w2v_model = Word2Vec(list_of_sent[100000:120000],min_count=5,size_
    ↪=100,workers=4)
w2v_words = list(w2v_model.wv.vocab)
for sent in list_of_sent[100000:120000]:
    sent_vec = np.zeros(100)
    count_words = 0
    for words in sent:
        if words in w2v_words:
            vec = w2v_model.wv[words]
            sent_vec += vec
            count_words +=1
    if count_words !=0:
        sent_vec = sent_vec/count_words
    X_cv_avg_w2v.append(sent_vec)

```

/home/niranjana/anaconda3/lib/python3.7/site-packages/gensim/models/base\_any2vec.py:743: UserWarning: C extension not loaded, training will be slow. Install a C compiler and reinstall gensim for fast training.

"C extension not loaded, training will be slow. "

```

[35]: X_test_avg_w2v = []
w2v_model = Word2Vec(list_of_sent[120000:140000],min_count=5,size_
    ↪=100,workers=4)
w2v_words = list(w2v_model.wv.vocab)
for sent in list_of_sent[120000:140000]:
    sent_vec = np.zeros(100)
    count_words = 0
    for words in sent:
        if words in w2v_words:
            vec = w2v_model.wv[words]
            sent_vec += vec
            count_words +=1
    if count_words !=0:

```

```
sent_vec = sent_vec/count_words
X_test_avg_w2v.append(sent_vec)
```

/home/niranjana/anaconda3/lib/python3.7/site-packages/gensim/models/base\_any2vec.py:743: UserWarning: C extension not loaded, training will be slow. Install a C compiler and reinstall gensim for fast training.

"C extension not loaded, training will be slow. "

```
[36]: import pickle
pickle_in1 = open("/home/niranjana/Downloads/AppliedAI/X_train_avg_w2v", "wb")
pickle_in2 = open("/home/niranjana/Downloads/AppliedAI/X_cv_avg_w2v", "wb")
pickle_in3 = open("/home/niranjana/Downloads/AppliedAI/X_test_avg_w2v", "wb")

pickle.dump(X_train_avg_w2v, pickle_in1)
pickle.dump(X_cv_avg_w2v, pickle_in2)
pickle.dump(X_test_avg_w2v, pickle_in3)
pickle_in1.close()
pickle_in2.close()
pickle_in3.close()
```

```
[1]: import pickle
pickle_out1 = open("/home/niranjana/Downloads/AppliedAI/X_train_avg_w2v", "rb")
pickle_out2 = open("/home/niranjana/Downloads/AppliedAI/X_cv_avg_w2v", "rb")
pickle_out3 = open("/home/niranjana/Downloads/AppliedAI/X_test_avg_w2v", "rb")

X_train_avg_w2v = pickle.load(pickle_out1)
X_cv_avg_w2v = pickle.load(pickle_out2)
X_test_avg_w2v = pickle.load(pickle_out3)
pickle_out1.close()
pickle_out2.close()
pickle_out3.close()
```

```
[15]: myList = list(range(0,50))

#creating odd list for k in Knn
neighbor = [x for x in myList if x%2 != 0]

cv_scores = []
for k in neighbor:
    classifier = KNeighborsClassifier(n_neighbors=k, algorithm='brute', n_jobs=1)
    classifier.fit(X_train_avg_w2v, y_train[0:100000])
    for i in range(0, len(X_cv_avg_w2v), 1000):
        y_pred = []
        y_pred.append(classifier.predict(X_cv_avg_w2v[i:i+1000]))
        y_pred_flat = [item for sublist in y_pred for item in sublist]
```

```

cv_scores.append(f1_score(y_cv[i:i+1000],y_pred_flat,average='micro'))

#Determining optimal value of k

optimal_k_brute = neighbor[cv_scores.index(max(cv_scores))]
print("optimal number of k is: ", optimal_k_brute)


plt.plot(neighbor,cv_scores)
for xy in zip(neighbor, np.round(cv_scores,3)):
    plt.annotate('%s, %s' % xy, xy=xy, textcoords='data')

plt.xlabel('Number of Neighbors K')
plt.ylabel('f1_score')
plt.show()

# print("f1_score for each k value is : ", np.round(cv_scores,3))

```

optimal number of k is: 5



```

[38]: y_pred = []
      clf = KNeighborsClassifier(n_neighbors=optimal_k_brute)
      clf.fit(X_train_avg_w2v,y_train[0:100000])
      for i in range(0,len(X_test_avg_w2v),1000):
          y_pred.append(clf.predict(X_test_avg_w2v[i:i+1000]))
      # evaluate accuracy
      y_pred_flat = [item for sublist in y_pred for item in sublist]

```


```

acc = f1_score(y_test[0:20000], y_pred_flat, average='micro')*100
print('\nThe accuracy of the knn classifier for k = %d is %f%%' % (
    optimal_k_brute, acc))
fpr, tpr, thresholds = metrics.roc_curve(y_test[0:20000], y_pred_flat,
    pos_label=2)
df = pd.DataFrame(confusion_matrix(y_test[0:20000], y_pred_flat))
sns.heatmap(df, annot=True, fmt="d")
plt.show()

```

The accuracy of the knn classifier for k = 5 is 82.965000%

/home/niranjana/anaconda3/lib/python3.7/site-packages/sklearn/metrics/ranking.py:571: UndefinedMetricWarning: No positive samples in y\_true, true positive value should be meaningless  
UndefinedMetricWarning)



```

[41]: y_pred_proba = []
      for i in range(0, len(X_test_avg_w2v), 1000):
          y_pred_proba.append(clf.predict_proba(X_test_avg_w2v[i:i+1000]))

```

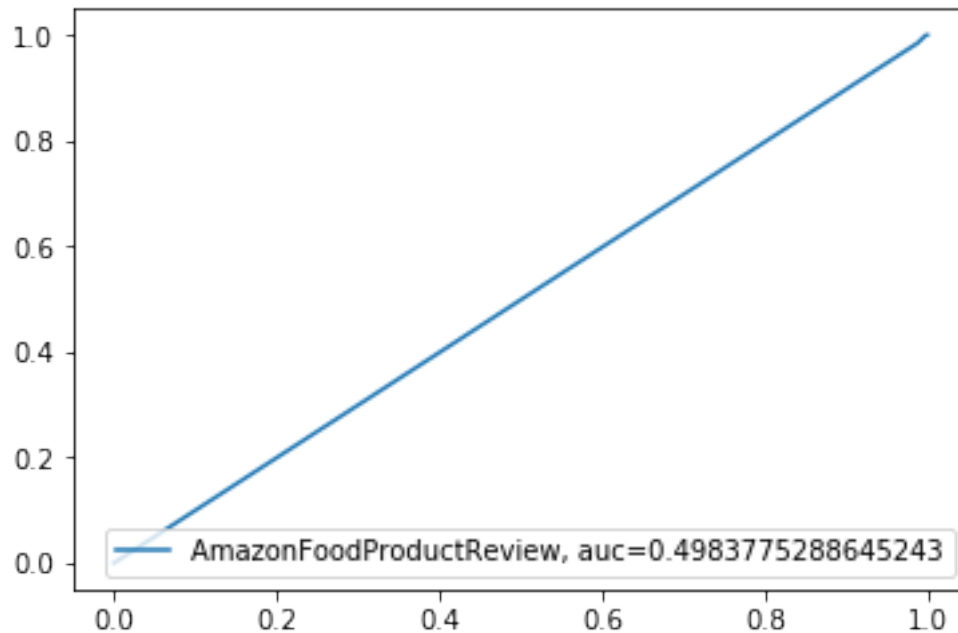
```

[42]: y_pred = []
      for i in y_pred_proba:
          for j in i:
              y_pred.append(j[1])

```



```
[43]: fpr, tpr, thresholds = metrics.roc_curve(y_test[0:20000], y_pred)
auc = metrics.roc_auc_score(y_test[0:20000], y_pred)
plt.plot(fpr,tpr,label="AmazonFoodProductReview, auc="+str(auc))
plt.legend(loc=4)
plt.show()
```



```
[44]: myList = list(range(0,50))

#creating odd list for k in Knn
neighbor = [x for x in myList if x%2 != 0]

cv_scores = []
for k in neighbor:
    classifier = KNeighborsClassifier(n_neighbors=k,algorithm='kd_tree',n_jobs=1)
    classifier.fit(X_train_avg_w2v,y_train[0:100000])
    for i in range(0,len(X_cv_avg_w2v),1000):
        y_pred = []
        y_pred.append(classifier.predict(X_cv_avg_w2v[i:i+1000]))
        y_pred_flat = [item for sublist in y_pred for item in sublist]
        cv_scores.append(f1_score(y_cv[i:i+1000],y_pred_flat,average='micro'))

#Determining optimal value of k

optimal_k_kd_tree = neighbor[cv_scores.index(max(cv_scores))]
```

```

print("optimal number of k is: ", optimal_k_kd_tree)


plt.plot(neighbor,cv_scores)
for xy in zip(neighbor, np.round(cv_scores,3)):
    plt.annotate('%s, %s' % xy, xy=xy, textcoords='data')

plt.xlabel('Number of Neighbors K')
plt.ylabel('f1_score')
plt.show()

# print("f1_score for each k value is : ", np.round(cv_scores,3))

```

optimal number of k is: 5



```


[45]: y_pred = []
      clf = KNeighborsClassifier(n_neighbors=optimal_k_kd_tree)
      clf.fit(X_train_avg_w2v,y_train[0:100000])
      for i in range(0,len(X_test_avg_w2v),1000):
          y_pred.append(clf.predict(X_test_avg_w2v[i:i+1000]))
      # evaluate accuracy
      y_pred_flat = [item for sublist in y_pred for item in sublist]
      acc = f1_score(y_test[0:20000], y_pred_flat,average='micro')*100
      print('\nThe accuracy of the knn classifier for k = %d is %f%%' %
            ↪(optimal_k_kd_tree, acc))
      fpr, tpr, thresholds = metrics.roc_curve(y_test[0:20000], y_pred_flat,
            ↪pos_label=2)

```

```
df = pd.DataFrame(confusion_matrix(y_test[0:20000],y_pred_flat))
sns.heatmap(df,annot=True,fmt="d")
plt.show()
```

The accuracy of the knn classifier for k = 5 is 82.965000%

```
/home/niranjana/anaconda3/lib/python3.7/site-
packages/sklearn/metrics/ranking.py:571: UndefinedMetricWarning: No positive
samples in y_true, true positive value should be meaningless
UndefinedMetricWarning)
```



```
[46]: y_pred_proba = []
for i in range(0,len(X_test_avg_w2v),1000):
    y_pred_proba.append(clf.predict_proba(X_test_avg_w2v[i:i+1000]))
```

```
[47]: y_pred = []
for i in y_pred_proba:
    for j in i:
        y_pred.append(j[1])
```

```
[48]: fpr, tpr, thresholds = metrics.roc_curve(y_test[0:20000], y_pred)
auc = metrics.roc_auc_score(y_test[0:20000], y_pred)
plt.plot(fpr,tpr,label="AmazonFoodProductReview, auc="+str(auc))
plt.legend(loc=4)
plt.show()
```



#### tfidf-Word2vec as vectorizer

```
[17]: # TF-IDF weighted Word2Vec
vec = TfidfVectorizer()
vec.fit_transform(X_train,y_train)
tfidf_feat = vec.get_feature_names() # tfidf words/col-names
dictionary = dict(zip(vec.get_feature_names(), list(vec.idf_)))
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val =
→tfidf
tfidf_train_vectors = []; # the tfidf-w2v for each sentence/review is stored in
→this list
row=0;
w2v_model = Word2Vec(list_of_sent[0:100000],min_count=5,size=100,workers=4)
w2v_words = list(w2v_model.wv.vocab)
for sent in tqdm(list_of_sent[0:100000]): # for each review/sentence
    sent_vec = np.zeros(100) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            if word in tfidf_feat:
                vect = w2v_model.wv[word]
#                tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
# to reduce the computation we are
# dictionary[word] = idf value of word in whole corpus
# sent.count(word) = tf value of word in this review
                tf_idf = dictionary[word]*(sent.count(word)/len(sent))
```

```

        sent_vec += (vect * tf_idf)
        weight_sum += tf_idf
    else:
        break
if weight_sum != 0:
    sent_vec /= weight_sum
tfidf_train_vectors.append(sent_vec)
row += 1

```

/home/niranjana/anaconda3/lib/python3.7/site-packages/gensim/models/base\_any2vec.py:743: UserWarning: C extension not loaded, training will be slow. Install a C compiler and reinstall gensim for fast training.

"C extension not loaded, training will be slow. "  
100%| | 100000/100000 [3:21:42<00:00, 6.19it/s]

```

[18]: # TF-IDF weighted Word2Vec
vec = TfidfVectorizer()
vec.fit_transform(X_train,y_train)

tfidf_feat = vec.get_feature_names() # tfidf words/col-names
dictionary = dict(zip(vec.get_feature_names(), list(vec.idf_)))
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val =
→tfidf
tfidf_cv_vectors = []; # the tfidf-w2v for each sentence/review is stored in
→this list
row=0;
w2v_model = Word2Vec(list_of_sent[100000:120000],min_count=5,size=100,workers=4)
w2v_words = list(w2v_model.wv.vocab)
for sent in tqdm(list_of_sent[100000:120000]): # for each review/sentence
    sent_vec = np.zeros(100) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            if word in tfidf_feat:
                vec = w2v_model.wv[word]
#                tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
# to reduce the computation we are
# dictionary[word] = idf value of word in whole corpus
# sent.count(word) = tf value of word in this review
                tf_idf = dictionary[word]*(sent.count(word)/len(sent))
                sent_vec += (vec * tf_idf)
                weight_sum += tf_idf
            else:
                break
    if weight_sum != 0:
        sent_vec /= weight_sum

```

```
tfidf_cv_vectors.append(sent_vec)
row += 1
```

/home/niranjana/anaconda3/lib/python3.7/site-packages/gensim/models/base\_any2vec.py:743: UserWarning: C extension not loaded, training will be slow. Install a C compiler and reinstall gensim for fast training.

"C extension not loaded, training will be slow. "  
100%| | 20000/20000 [39:48<00:00, 7.59it/s]

```
[19]: # TF-IDF weighted Word2Vec
vec = TfidfVectorizer()
vec.fit_transform(X_train,y_train)

tfidf_feat = vec.get_feature_names() # tfidf words/col-names
dictionary = dict(zip(vec.get_feature_names(), list(vec.idf_)))
# final_tf_idf is the sparse matrix with row= sentence, col=word and cell_val =
→tfidf
tfidf_test_vectors = []; # the tfidf-w2v for each sentence/review is stored in
→this list
row=0;
w2v_model = Word2Vec(list_of_sent[120000:140000],min_count=5,size=100,workers=4)
w2v_words = list(w2v_model.wv.vocab)
for sent in tqdm(list_of_sent[120000:140000]): # for each review/sentence
    sent_vec = np.zeros(100) # as word vectors are of zero length
    weight_sum =0; # num of words with a valid vector in the sentence/review
    for word in sent: # for each word in a review/sentence
        if word in w2v_words:
            if word in tfidf_feat:
                vec = w2v_model.wv[word]
#                tf_idf = tf_idf_matrix[row, tfidf_feat.index(word)]
# to reduce the computation we are
# dictionary[word] = idf value of word in whole corpus
# sent.count(word) = tf value of word in this review
                tf_idf = dictionary[word]*(sent.count(word)/len(sent))
                sent_vec += (vec * tf_idf)
                weight_sum += tf_idf
            else:
                break
    if weight_sum != 0:
        sent_vec /= weight_sum
    tfidf_test_vectors.append(sent_vec)
    row += 1
```


/home/niranjana/anaconda3/lib/python3.7/site-packages/gensim/models/base\_any2vec.py:743: UserWarning: C extension not loaded, training will be slow. Install a C compiler and reinstall gensim for fast

```
training.  
"C extension not loaded, training will be slow. "  
100%|      | 20000/20000 [42:57<00:00, 10.68it/s]
```

```
[21]: import pickle  
pickle_in1 = open("/home/niranjana/Downloads/AppliedAI/tfidf_train_vectors","wb")  
pickle_in2 = open("/home/niranjana/Downloads/AppliedAI/tfidf_cv_vectors","wb")  
pickle_in3 = open("/home/niranjana/Downloads/AppliedAI/tfidf_test_vectors","wb")  
  
pickle.dump(tfidf_train_vectors,pickle_in1)  
pickle.dump(tfidf_cv_vectors,pickle_in2)  
pickle.dump(tfidf_test_vectors,pickle_in3)  
pickle_in1.close()  
pickle_in2.close()  
pickle_in3.close()
```

```
[29]: myList = list(range(0,50))  
  
#creating odd list for k in Knn  
neighbor = [x for x in myList if x%2 != 0]  
  
cv_scores = []  
for k in neighbor:  
    classifier = KNeighborsClassifier(n_neighbors=k,algorithm='brute',n_jobs=1)  
    classifier.fit(tfidf_train_vectors,y_train[0:100000])  
    for i in range(0,len(tfidf_cv_vectors),1000):  
        y_pred = []  
        y_pred.append(classifier.predict(tfidf_cv_vectors[i:i+1000]))  
        y_pred_flat = [item for sublist in y_pred for item in sublist]  
        cv_scores.append(f1_score(y_cv[i:i+1000],y_pred_flat,average='micro'))  
  
#Determining optimal value of k  
  
optimal_k_tfidf_brute = neighbor[cv_scores.index(max(cv_scores))]  
print("optimal number of k is: ", optimal_k_tfidf_brute)  
  
plt.plot(neighbor,cv_scores)  
for xy in zip(neighbor, np.round(cv_scores,3)):  
    plt.annotate('(%s, %s)' % xy, xy=xy, textcoords='data')  
  
plt.xlabel('Number of Neighbors K')  
plt.ylabel('f1_score')  
plt.show()  
  
# print("f1_score for each k value is : ", np.round(cv_scores,3))
```

```
optimal number of k is:  11
```




```
[33]: y_pred = []
      clf = KNeighborsClassifier(n_neighbors=optimal_k_tfidf_brute)
      clf.fit(tfidf_train_vectors,y_train[0:200000])
      for i in range(0,len(tfidf_test_vectors),1000):
          y_pred.append(clf.predict(tfidf_test_vectors[i:i+1000]))
      # evaluate accuracy
      y_pred_flat = [item for sublist in y_pred for item in sublist]
      acc = f1_score(y_test[0:20000], y_pred_flat,average='micro')*100
      print('\nThe accuracy of the knn classifier for k = %d is %f%%' %_
            ↪(optimal_k_tfidf_brute, acc))
      fpr, tpr, thresholds = metrics.roc_curve(y_test[0:20000], y_pred_flat,_
            ↪pos_label=2)
      df = pd.DataFrame(confusion_matrix(y_test[0:20000],y_pred_flat))
      sns.heatmap(df,annot=True,fmt="d")
      plt.show()
```

The accuracy of the knn classifier for k = 11 is 82.875000%

```
/home/niranjana/anaconda3/lib/python3.7/site-
packages/sklearn/metrics/ranking.py:571: UndefinedMetricWarning: No positive
samples in y_true, true positive value should be meaningless
  UndefinedMetricWarning)
```






```
[37]: y_pred_proba = []
      for i in range(0, len(tfidf_test_vectors), 1000):
          y_pred_proba.append(clf.predict_proba(tfidf_test_vectors[i:i+1000]))
```

```
[38]: y_pred = []
      for i in y_pred_proba:
          for j in i:
              y_pred.append(j[1])
```

```
[39]: fpr, tpr, thresholds = metrics.roc_curve(y_test[0:20000], y_pred)
      auc = metrics.roc_auc_score(y_test[0:20000], y_pred)
      plt.plot(fpr, tpr, label="AmazonFoodProductReview, auc="+str(auc))
      plt.legend(loc=4)
      plt.show()
```



tfidf-Word2vec vectorizer with kd\_tree as algorithm

```
[36]: myList = list(range(0,50))

#creating odd list for k in Knn
neighbor = [x for x in myList if x%2 != 0]

cv_scores = []
for k in neighbor:
    classifier = KNeighborsClassifier(n_neighbors=k,algorithm='kd_tree',n_jobs=1)
    classifier.fit(tfidf_train_vectors,y_train[0:100000])
    for i in range(0,len(tfidf_test_vectors),1000):
        y_pred = []
        y_pred.append(classifier.predict(tfidf_test_vectors[i:i+1000]))
        y_pred_flat = [item for sublist in y_pred for item in sublist]
        cv_scores.append(f1_score(y_cv[i:i+1000],y_pred_flat,average='micro'))

#Determining optimal value of k


optimal_k_tfidf_kd_tree = neighbor[cv_scores.index(max(cv_scores))]
print("optimal number of k is: ", optimal_k_tfidf_kd_tree)

plt.plot(neighbor,cv_scores)
for xy in zip(neighbor, np.round(cv_scores,3)):
    plt.annotate('(%s, %s)' % xy, xy=xy, textcoords='data')
```


```
plt.xlabel('Number of Neighbors K')
plt.ylabel('f1_score')
plt.show()

# print("f1_score for each k value is : ", np.round(cv_scores,3))
```

optimal number of k is: 23



```
[40]: clf = KNeighborsClassifier(n_neighbors=optimal_k_tfidf_kd_tree)
clf.fit(tfidf_train_vectors,y_train[0:200000])
pred = clf.predict(tfidf_test_vectors)
# evaluate accuracy
acc = f1_score(pred,y_test[0:20000])*100
df = pd.DataFrame(confusion_matrix(pred,y_test[0:20000]))
sns.heatmap(df,annot=True,fmt="d")
plt.show()
print('\n\nThe accuracy of the knn classifier for k = %d is %f%%' %_
      ↪(optimal_k_tfidf_kd_tree, acc))
```



The accuracy of the knn classifier for k = 23 is 0.906611%

```
[2]: from prettytable import PrettyTable
x= PrettyTable()
x.field_names = ["Vectorizer","Model","Hyper parameter","Acurracy","Auc"]
x.add_row(["BOW","Brute","9","93.54",".6764"])
x.add_row(["BOW","kd_tree","9","93.54",".6764"])
x.add_row(["tf-idf","Brute","7","93.89",".79014"])
x.add_row(["tf-idf","kd_tree","7","93.89",".79014"])
x.add_row(["Avg-Word2Vec","Brute","13","94.21",".4983"])
x.add_row(["Avg-Word2Vec","kd_tree","5","82.9650",".4983"])
x.add_row(["tfidf-Word2Vec","Brute","11","82.8750",".5412"])
x.add_row(["tfidf-Word2Vec","kd_tree","23","90.66",".5412"])
print(x)
```

Vectorizer	Model	Hyper parameter	Acurracy	Auc
BOW	Brute	9	93.54	.6764
BOW	kd_tree	9	93.54	.6764
tf-idf	Brute	7	93.89	.79014
tf-idf	kd_tree	7	93.89	.79014
Avg-Word2Vec	Brute	13	94.21	.4983
Avg-Word2Vec	kd_tree	5	82.9650	.4983
tfidf-Word2Vec	Brute	11	82.8750	.5412
tfidf-Word2Vec	kd_tree	23	90.66	.5412

tfidf-Word2Vec	kd_tree	23		90.66		.5412	
+-----+-----+-----+-----+-----+							